Can agents talk about what they are doing? A proposal with Jason and speech acts

Valeria Seidita[a], Francesco Lanza[a], Angelo Maria Pio Sabella[a] and Antonio Chella[a,b]

[a]*Dipartimento di Ingegneria, Universitá degli Studi di Palermo, Italy*
[b]*ICAR-CNR National Research Council, Palermo, Italy*

Abstract

The dream of building robots and artificial agents that are more and more capable of thinking and acting like humans is growing by the day. Various models and architectures aim to mimic human behavior. In the course of our current research, we propose a solution to make actions and thought cycles of agents explainable by introducing inner speech into a multi-agent system. The reasons that led us to use inner speech as a self-modeling engine raised the definition of what inner speech is and how it affects cognitive systems. In this proposal, we used speech act to enable a coalition of agents to exhibit inner speech capabilities to explain their behavior, but also to guide and reinforce the creation of an inner model that is triggered by the decision-making process through actions applied to the surrounding world, but also to themselves. The BDI agent paradigm is used to keep agents rational and with the innate ability to act in a human-like manner. The proposed solution continues the research path that began with the definition of a cognitive model and architecture for human-robot teaming interaction, and aims to integrate the credible interaction paradigm into it.

Keywords

Human-Agent Interaction, Transparency, Jason, Inner Speech

1. Introduction

"What does a person think about before taking an action to achieve a goal?" — This question must be asked if one wants to reproduce the decision-making abilities proper of humans in an agent. Making a decision profitably is a very complex process that depends on many factors, especially when dealing with collaborative tasks and a highly variable environment. Getting an agent to make a decision in a fruitful way that takes into account capabilities such as independently choosing the best action or updating one's knowledge in pursuit specific goals or observing the behavior of others therefore presents a number of challenges.

When we talk about collaborative tasks, we inevitably refer to human-agent (or human-machine, or even human-robot) interaction. In a collaborative task, humans and agents cooperate and work together to achieve a common goal. The goal is achieved through communication and interaction between agents, agents and humans, and agents and the environment. For example, consider the behavior and interaction processes of the members of a team of humans. The goal

WOA 2022: Workshop "From Objects to Agents", September 1–3, 2021, Genoa, Italy
valeria.seidita@unipa.it (V. Seidita); francesco.lanza@unipa.it (F. Lanza); angelomariapio.sabella@community.unipa.it (A. M. P. Sabella); antonio.chella@unipa.it (A. Chella)
0000-0002-0601-6914 (V. Seidita); 0000-0003-4382-6366 (F. Lanza); 0000-0002-4325-759X (A. Chella)

is shared, each team member knows the goal, a set of actions (or tasks, i.e., one or more plans) to achieve it, his or her own capabilities, probably those of others, and all the known elements of the environment. In complex interactions and tasks, often not everything is known about the environment; not everyone knows everything, but can share knowledge. In addition, each team member's interaction with the environment can cause a change in the environment that can affect another member's actions.

So, when a team member decides to perform an action, he checks if this action is within his capabilities and if no other team member has already performed it or is performing it. If so, he can perform the action, delegate it to another team member, or not perform it and change it. Communication between team members and interaction with the environment are continuous and necessary for gaining new knowledge about oneself and the environment (other team members and actual objects), and they are the main source for changing the environment and increasing one's knowledge about it.

Even before deciding whether and how to take an action, there is a process of deciding or selecting the most appropriate action to achieve a goal. If the goal were easy to achieve, with a few actions already determined, and in an unchanging environment, humans would not even have to think and would slavishly and reactively perform the specified actions to achieve the goal. We are aware that such a simple, let us say atomic, situation for achieving complex goals is not common in teams.

There are many elements that underlie the decision-making process, including knowledge of one's own capabilities, continuous replanning when a goal is not achieved, and interaction with peers to gather information to gain new knowledge. In addition, in the context of collaboration and interaction, factors, such as the trust one team member has in another, his or her mental states (the inclination or willingness to perform an action), and the knowledge derived from the other's ability to explain the reasons for his or her actions or outcomes, must also be considered.

The interaction between humans and agents raises issues that are still being explored from a design process perspective. If everything about a system is not known a priori, and it evolves and changes during its execution phase, then it is not possible to define the properties of the system at design time and to study the requirements for the system completely. We have been working on this part of human-agent system development for years, and in this paper we also identify some requirements for human-agent interaction systems.

Our idea is based on the creation of a cognitive model and the corresponding agent architecture, whose modules allow structuring the agent's decision-making process, taking into account its internal states. Thus, we have combined aspects of self-modeling and Theory of Mind to integrate the elements of interaction already mentioned. On the implementation level, on the other hand, we have harnessed the power of the BDI agent paradigm [1, 2, 3] and one of the best known and most widely used agent languages, *Jason* [4, 5].

In the present work, we go a step further and make a very early step for extending what we have done so far to give the agent the ability to reason aloud, i.e., to externalize its reasoning and thus make its actions comprehensible to humans.

Our proposal is based on the recognition that humans are able to talk to themselves, either aloud or in a quiet voice, in order to improve their understanding of what they are doing, to self-regulate their behavior, and to increase their knowledge of what surrounds them by using the feedback that inner speech gives them. We propose to combine the concept of inner speech

with the concept of speech act to obtain a tool that allows us to move from the cognitive model to the agent architecture and then to the actual implementation. Moreover, this idea provides the input to explore the possibility of modifying the selection functions of the reasoning cycle of the *Jason* interpreter[4]. Thus for the cognitive model, we then propose a first validation through a simple example in which we have implemented inner speech through speech acts.

The rest of the paper is organized as follows: in section 2, we describe the work done so far in this context and then, in sections 3 and 4, we give an overview of the concepts of inner speech and speech acts with the aim of highlighting their characteristics and the motivation for using them; in sections 5 and 6, we provide an analysis of the requirements of the systems under study and of the model that is proposed; in section 7 the validation scenario of the model with its code is given; and finally, in section 8 some conclusions are drawn.

2. Decision Process in Human-Agent Interaction

Adaptivity, autonomy, and proactivity are the qualities or abilities that every human being must possess in order to be able to decide at any time what actions to take in order to achieve the desired results in the achievement of a goal. Whether it is humans or agents (or robots), someone must have the ability to make appropriate decision-making, and if the context is a human-agent interaction, the agent must have the ability to select an action and decide whether to perform it by itself or delegate it to someone else.

In [6] we proposed an architecture for modeling a human-robot team system. In that paper, we refer to a robot, but the idea starts from and can be applied to the concept of agents and multi-agent systems deployed in a robotic system. The proposed architecture uses the elements underlying human-agent interaction. First of all, it is based on the standard model of the mind [7], which considers the classical MAPE (Monitoring, Analyse, Process, Execute) process in the thought cycle of any intelligent system [8]. From this, we added the modules for representing autonomous and adaptive interactions: in particular, we considered that the inputs to the reasoning process, in addition to the goal, also come from the subdivision of the main goal into subgoals and from what we called motivations, the heart of the decision process and the elements that initiate it. Here are all the information and processes that represent the inner world of an agent, his beliefs, desires, intentions, knowledge, and skills, but also norms, rules, emotions, the degree of trust in the other, and everything else that can serve as input for action (see Figure 1).

A fundamental premise of this work is that we consider as the environment not only the external world of an agent (i.e., the objects on which it can perform actions) but also the internal (the self-model) and other agents. In this way, from a theoretical point of view, it is possible to model situations such as the justification of one's actions or, more importantly for decision making and replanning, the anticipation of actions. Anticipation is the process that allows generating a "current situation" i.e., the state of the world, from a selected possible action. It allows imagining the outcome of an action and actually implementing it only if its outcome is consistent with the post-conditions of the selected goal. In [9, 10] we proposed a prototype for part of the described architecture and implementation using the BDI paradigm and *Jason* language, respectively.

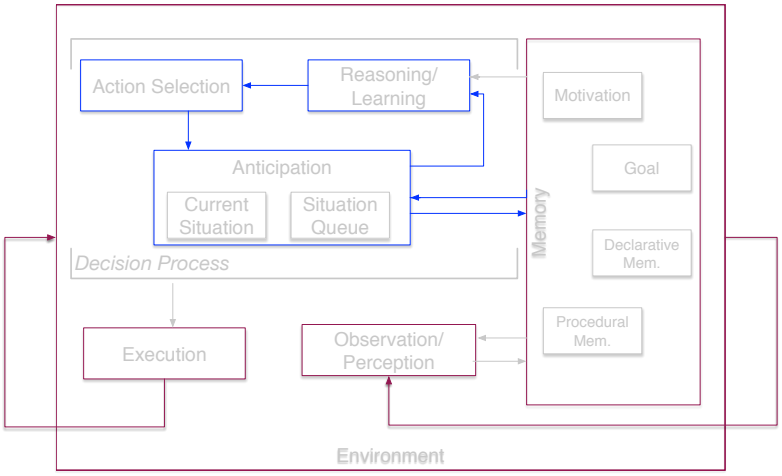

Figure 1: The abstract architecture for creating human-agent (robot) interaction based on human-human interaction in a team working together in the same rapidly changing environment to achieve a known and common goal.

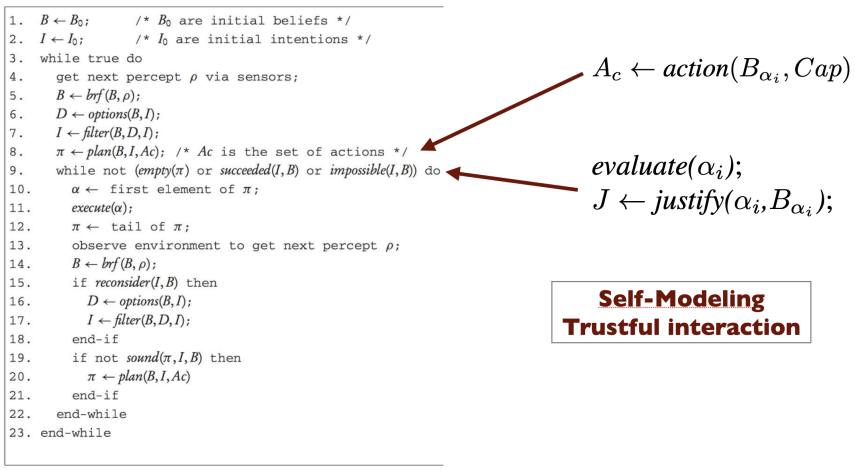

Figure 2: The algorithm for agent's practical reasoning and the extension for adding self-modeling and trustful interactions [4, 11].

In the first work, we paid special attention to the module of anticipation and robot knowledge enhancement. Basically, we developed a tool to look into the robot's mind while it is performing its collaborative tasks. In the second work, we proposed a variation of the reasoning loop that lies below the *Jason* interpreter to incorporate general motivation into the language and its management in the same way that *Jason* does in managing beliefs.

Finally, in [11] we used the trust model developed by Falcone and Castelfranchi [12, 13] in conjunction with the practical reasoning theory to build a model for the robot's decision. The idea was to extend the deliberative process and the belief base representation to allow the robot to decompose a plan into a series of actions. Each action is directly linked to the knowledge

useful for its execution and is computed by a function

$$J \leftarrow justify(\alpha_i, B_{\alpha_i}) \tag{1}$$

In this way, the robot creates and maintains a model of "itself" and can justify the results of its actions. Justification is a key outcome of the application of self-modeling capabilities, as well as a useful means of improving trust interactions. It takes place in the execution phase, i.e., at the beginning of the justification cycle, a *Jason* agent updates all its beliefs and intentions, determines its desires, and selects some of them that become intentions to which it assigns a plan (line 1 to line 8 in Figure 2). After that, the agent usually processes the stack of actions to decide which ones to execute. At this point, we have added a new function

$$A_c \leftarrow action(B_{\alpha_i}, Cap) \tag{2}$$

that allows us to associate part of the belief base with the capabilities of the robot, which in this way is able to justify its actions through the function $\mathcal{J}$.

The idea we are now proposing is to extend this even further and give the agent the ability to activate what is called inner speech, the ability to speak to oneself by making oneself the object of thought. The implementation is realizing using the instrument of speech acts, the basic theory of which is perfectly consistent with the psychological theory of inner speech. To our knowledge, there is no approach in the literature, especially in the field of agents and robotics, that combines inner speech with speech acts both conceptually and in terms of implementation.

3. Inner Speech

The term inner speech has several connotations in the literature. First, it is a concept developed and used in psychology. It is often associated with the concept of self-awareness and self-consciousness [14, 15]. Self-awareness is "the ability to become the object of one's attention" [16] and takes into account three possible elements: the social environment, the physical world, and the self. Self-awareness also includes the ability to direct attention to one's mental state, i.e., perceptions, feelings, attitudes, intentions, emotions, etc. Morin was the first in [17] to link the concept of inner speech to the ability to think about oneself.

Inner speech is a concept used mainly in cognitive development and executive function theory. It is a way of reflecting on one's own experiences. People generally reflect on their own experiences in different ways [18]. Inner speech plays an important role in self-regulation of cognition and behavior, so it is used in people for both data collection and behavior regulation, and is also considered a motivational tool.

A definition of inner speech is very complex. Some authors define inner speech as the subjective experience of language in relation to one's actions, feelings, and experiences. It is often used as a synonym for thinking. Most authors in the literature state that it is better to use the term mental process instead. In a number of seminal studies, inner speech is associated with cognition and behavior or even to rehearsal and working memory. In each case, inner speech is recognized as a feature of the developmental process. Thus, it provides the output to activate a developmental process. According to Morin, inner speech is the activity of silent self-talk. There

are many other synonyms or equivalents, with some differences between adults and children, for example, self-talk, internal dialog, private speech or egocentric speech, or self-verbalization. The latter two activities generally refer to the behavior of children who comment on their own actions without caring whether they are understood or not. What interests us from a technical point of view is that inner speech, as has been shown, serves self-direction, self-regulation, problem solving, planning, and memory. Another important point we want to keep in mind is that inner speech serves information retrieval. Or rather, self-information, as it is called by Morin and Everett. As mentioned earlier, a person who needs to perform an action may talk to himself. This conversation serves him to identify data and processes. So to take actions and on what elements. We claim in this moment of the thinking process one can associate the possibility or the abilities, in the case of an agent, to externalize his thinking and thus make it explainable or transparent.

The sources of inner speech are the social milieu and the physical environment. For our purposes, however, we need to consider only the latter. Indeed, the physical environment contains a set of stimuli that enable a person, in our case an agent, to reason about the stimuli emanating from the external environment and to make appropriate decisions. In interaction with people, the use of inner speech or loud inner speech can be a way to make cognitive and decision-making processes transparent. And it is also a way to perform verbal mediation with oneself to support certain activities. Often, especially with children, inner speech is seen or used as an accompaniment or constant commentary on what they are doing. This is exactly what we want to do. In this paper, we take our cue from [19] and propose to give an agent the ability to comment aloud on what it is doing. The actions it chooses to take, how it acts. In this way, we can make an agent's decision-making processes and actions transparent. Vygotsky in [20] also says that verbal accompaniment of action is much more pronounced in tasks that present extreme difficulties. Since an agent has to make decisions for which it does not have a precise plan when it is in a difficult situation, modeling and implementing the inner speech in an agent can lead to more efficient decision making process Our idea is to build on a parallel to inner language in development process of children. In our work, we are mainly concerned with the development and design of robots that interact autonomously with humans. The robot, working in an environment that is not fully know, is in some ways like a developing child.

4. Speech acts

In this section, we provide a brief overview of speech acts and focus on the elements that made us curious about the possibility of using them to implement the inner speech module.

Following our research topics, we are focused on the definition of techniques useful for enabling BDI agents to show inner speech capabilities for increasing the level of trust during human-robot interaction.

In short words, inner speech may be considered as the capability of an artificial agent to speak loudly with itself. The multi-agent paradigm provides a set of functionalities that enables agents to speak with other agents. Communication between agents changes the state of the world of every single agent that receives a new message. In general, this behavior may be a matter of the speech act theory proposed by Searle [21] and Austin[22].

Performative	Description
tell	*sender* intends *receiver* to believe (that sender believes) the literal in the message's content to be true;
untell	*sender* intends *receiver* not to believe (that sender believes) the literal in the message's content to be true;
achieve	*sender* requests *receiver* to try and achieve a state of affairs where the literal in the message content is true (i.e. sender is delegating a goal to receiver);
unachieve	*sender* requests *receiver* to drop the goal of achieving a state of affairs where the message content is true;
askOne	*sender* wants to know if the content of the message is true for *receiver* (i.e. if there is an answer that makes the content a logical consequence of receiver's belief base, by appropriate substitution of variables);
askAll	*sender* wants all of *receiver*'s answers to a question;
tellHow	*sender* informs *receiver* of a plan (sender's know-how);
untellHow	*sender* requests that *receiver* disregard a certain plan (i.e. delete that plan from its plan library);
askHow	*sender* wants all of *receiver*'s plans that are relevant for the triggering event in the message content.

Table 1
This table contains a list of all available performative in *Jason*. Every performative's description was taken faithfully from [23].

Bordini et al. [23] described the speech act and how it works in an agent communication module. The basic principle of the speech act theory lies in the meaning of *language*. The principle can be summarized by assuming language as action. An artificial agent uses utterance to inform other agents about changes or to exchange novel information that concerns the surrounding world. The utterance produced by an agent changes, effectively, the state of the world in terms of beliefs, desires, and intentions owned by a hearer agent. In the Austin's speech act framework [22], speech act are *locutionary*, *illocutionary* and, *perlocutionary*. *Locutionary acts* concern what was said and meant, *illocutionary acts* represent what was done and, *perlocutionary acts* describe what happened as a result. According with this theory, Searle [21] identified various types of speech act: (*i*) representatives, (*ii*) directives, (*iii*) commisives, (*iv*) expressives and, (*v*) declarations. Theoretically, speech acts are composed by two parts: (*i*) a performative verb, and (*ii*) a propositional content.

A multi-agent system supports basic speech act in its communication module through the adoption of agent communication languages [24], such as the *Knowledge Query and Manipulation Language* (KQML) [25, 26] or FIPA [27, 28, 29]. For instance, as denoted in [23], speech acts in *Jason* may be used though the usage of a set of internal action that let agents communicate one with other.

Accordingly with Bordini [23], the BDI framework *Jason* is able to perform illocutionary speech act, also know as performative speech act. This is possible thanks to the set of internal actions owned by each agent. Looking at the reasoning cycle of the *Jason* agent, the reasoning

cycle starts even checking mails. In fact, each time that an agent receives a message, this message is composed by:

<sender, illoc_force, content>

where sender is the name usage by the sender agent in the multi-agent system, illoc_force denotes the intention of the sender, and the content is a term that depends on the type of illocutionary force, and it represents the content of the message. In *Jason*, a basic illocutionary speech act may be send from a sender agent to a receiver using the *.send* internal action. Depending on what the agent wants to communicate, it can use different types of illocutionary speech act or performative. *Jason* provides a set of semantic performative to let an agent communicate. The *performatives*' name follow the KQML standard [30, 31]. In Table 1 a list of performatives, every definition was taken faithfully from [23].

Given the list of all possible illocutionary speech acts, then, the point is to exploit them in order to let the agent speak to himself, to show the capacity for inner speech. In this way, inner speech lets the agent's mental state change, as well as the agent's intention and behavior. Moreover, inner speech, by definition, makes the agent's decision-making process explainable.

5. Requirements of a trustful and explainable human-agent software

From an engineering perspective, the primary elements to consider when designing a multi-agent system, or at least an intelligent system that can autonomously interact and collaborate with a human, are the following: the goals of the agent or team (i.e., the agent system in our case), the behaviors and thus the tasks and actions that an agent must perform, and the communication, message exchange between agents to achieve the common task.

At design time, a designer must then design a knowledge base that represents the world in which the agents operate, the external world, and based on what was discussed earlier, for each agent, the robot's self, i.e., knowledge about itself as it is, which is part of the working environment anyway. For the question of what the software needs to show when it is running, we can go back to the scenario we used in our previous work [32, 33, 34, 11]. Let us say a team of humans and robots needs to set a table, following some rules of etiquette. For example, knives are placed to the right of the plate, forks to the left, etc. The robots can be designed using a multi-agent system that is given all the *plans* for setting a table, that is, all the possible *actions*, all the elements of the table, and also all the *rules* to be followed.

In this scenario, for simplicity, we do not assume a mutable environment, i.e., there are no objects or rules that are added at runtime, and there are no external events that could interfere with the activation of the plans. Also, we consider one agent deployed in each robot. At the beginning of the collaboration, each team element selects the action it wants to perform from those specified in the plan. Suppose one of the robots has committed to picking up and putting down a fork. However, the robot sees that another robot or a human has already placed the fork on the table. It interrupts the action and activates the internal speech module to explain to itself and others why it interrupted the action. It contextually changes its knowledge base about the presence and position of the fork. The robot also recognizes that the fork is placed

in the wrong way or in the wrong place. A rule of etiquette has been broken, and again the robot activates its inner speech to communicate externally that the other team member has not performed his task correctly.

The full design and development of this example scenario includes many other cases, such as one in which the robot fails to complete an action and reasons to itself why this happened. In each case, the robot interacts with other elements of the team to achieve a goal, but also to explain why it does what it does and to have the elements with which to make decisions and plan at runtime. The first element we know underlies the agent design paradigm, while the second allows us to make an agent explainable and trustworthy at the same time.

From the in-depth analysis of the previous scenario and other similar scenarios we have worked on, we have identified a list of requirements that the agent system must expose at runtime in order to interact with humans in a trustworthy and explainable way.

- **Selecting an action** - the agent must be able to select the appropriate action based on a defined goal and decide whether to carry it out or delegate it to another team member, justifying his decision;
- **Capturing and verifying preconditions** - the agent must constantly identify and evaluate elements of the environment, observe other team members, and associate with each action the correct condition (state of the environment) to perform the action. In this case, part of the agent's knowledge may relate to the norms governing the work environment, but also to his own internal state and the mental state of himself and other members;
- **Updating the knowledge base** - the system must be able to update its knowledge base at runtime. Although we have envisioned an immutable environment in this work, it must be taken into account that any interaction of an agent with its peers or with the environment changes both external and internal states. An example would be the fork mentioned above that changed its position, or, in a broader and more complex scenario, the acquisition of knowledge that a team member can or will only perform actions under certain conditions, or even the acquisition of knowledge about how to achieve a goal through an action or plan that is not specified by the design;
- **Review and reasoning about the results of an action** - after each action, the agent must review his inner state and that of the world and decide whether the action was successful; this must be done both for his own actions and for the actions of others. He must then be able to reason and justify why an action was not successful, if necessary;
- **Considerations of team behavioral norms** (as a function of the shared objective of the team) - behavioral norms refer to all those constraints in the work environment that influence or constrain the agent's behavior in some way. As with world objects or inner world, they are part of the knowledge base of the agent system, given in an initial formulation in the design phase. Norms may then need to be updated in the execution phase, at runtime. In the scenario we illustrated above, the standards could be the etiquette rules for setting a table, and in our case they do not change. In either case, the agent must be able to reason and activate an internal discourse about the norms that constrain the system in order to justify its actions, explain its own behavior and that of others, or even add preconditions for certain actions.

As indicated in the introduction, it is necessary to consider the system both from the point of view of all that can be planned at design time and from the point of view of what is to be handled by the agents at runtime. Planning when the designer cannot specify all the plans, reasoning about the results of actions, and updating knowledge. The tools currently available for the design and development of agent-based systems, in particular BDI paradigm and *Jason* language, provide valuable and proven help for the requirements we have for design time and some for runtime. In the next sections, we will illustrate the agent approach we used to develop the part dealing with credibility, trust, and explainability, using a model that allows us to endow agents with the ability to execute and then reveal an inner speech.

6. Agents are able to explain themselves: the proposed approach

In our earlier work, we deepened and developed the concept of justification, which had more to do with establishing a degree of confidence in the agent's abilities, which led us to focus solely on the action and its outcome. Thus, the agent or robot justifies why it did not perform an action or chose to delegate it (or not commit it) by computing the self-model. In this work, we extend the concept of justification on the outcome of the action to the entire action, from selection to execution, and use the results of reasoning to update the agent's knowledge base and explain to the outside world what it is doing.

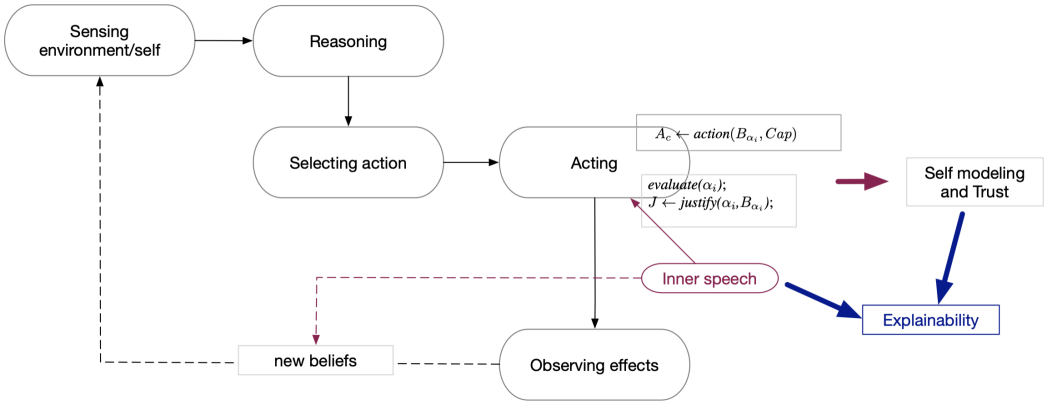

Figure 3: Key ideas for implementing inner speech in agents.

In Figure 3 we show our proposed behavioral view of the modules in the architecture shown the Figure 1. We did not include a start and end point because we want to represent the continuous reasoning cycle of an agent. An agent reasons basing on input from the environment, both external and internal, and selects an action from the set provided by the designer (in the case of a human-agent interaction, the agent may also ask the human for suggestions on what to do, adding actions and plans to its knowledge base, but this is outside the scope of the work

proposed here). Once the action is selected, the two functions mentioned in section 2 are activated to create a rationale for the actions. Typically, the agent then observes the effects of its actions and the resulting changes in the external world, and updates its beliefs through perception. We also added and account for the updating of beliefs that are part of the internal world, so that at runtime we create and update the agent's self. We also add the inner speech module at the point where we use the $\mathcal{J}$ function. Before and after the execution of each action, the agent starts its inner speech and generates new beliefs, in addition to those generated by observing the effects on the environment.

From a technological point of view, the BDI agent paradigm and the *Jason* interpreter reasoning cycle are well suited for developing the parts we focused on in Figure 3. At runtime, a BDI agent consists of the Belief set, the Plan set, the Intention set, the Event set, and Action as well as the selection functions managed by the interpreter: S_E, S_O, and S_I [4].

Events can be external and internal. External events are generated by the perception of the environment and correspond to the deletion or addition of beliefs. Internal events are generated by the agent when executing a plan and they do not affect the environment. Thus, events are the ones most closely associated with the execution phase. Some examples are the addition of achievement or test goals.

Our idea is to handle internal events for linking them with speech act to generate an explanation for agents' actions.

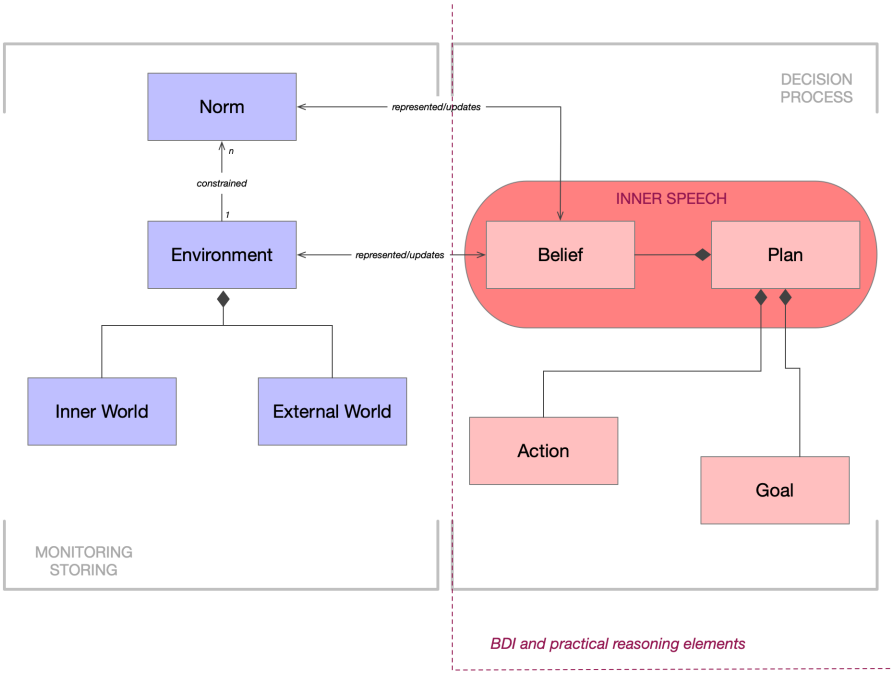

Figure 4: Design abstraction for human-agent interaction and inner speech.

The plan represents how the agent should act to achieve a certain goal, taking into account the belief, the action to be performed and the goal to be achieved. Plans, then, are the elements that make inner discourse possible, and it is on these that we now focus.

The idea we had is simple from the theoretical point of view and finds its natural implementation in speech acts. In an agent system, communication is based on the theory of speech acts [21], briefly described in the section 4. Since inner speech is a way of thinking about oneself, we claim that under certain conditions a speech act can address the agent himself. We assume that each agent in each plan, in addition to taking actions to achieve team goals, has the possibility of sending a message to itself aimed at changing its beliefs or goals, while making this message available to the outside world. In the work we propose in this paper, the communication with the outside world is done through the console; in the near future, a real dialogue between agent and human will be implemented.

The elements in Figure 4 show what is described above and the rationale for translating towards the inner speech implementation phase through *Jason* agents.

The operating environment, whether internal or external, can be represented by the agent's beliefs. The environment is governed by norms, which are also represented by beliefs. Beliefs are the main elements that make up a plan, along with actions and goals. The left part of the figure shows the elements that allow us to instantiate the monitoring and storage modules of the abstract architecture. On the right side, on the other hand, we can instantiate the reasoning module, which is based on the logical foundation of the BDI paradigm and practical reasoning. If we had used a different paradigm, these elements would look different, but the logical basis would not change. The advantage we have gained from our choice is that we can easily connect the theoretical model with its practical implementation by using the *Jason* interpreter's reasoning cycle.

Through beliefs, we can link actions to the environment and thus have or maintain knowledge about what actions can or cannot be performed and on what objects this can occur.

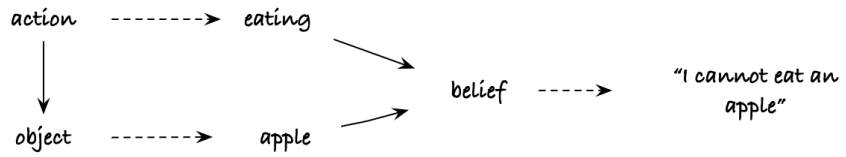

For example, if we have an object apple, an action eating, and the relation between them realizes a norm of "not", then we can have the agent infer that it cannot eat the apple. If we make the representation more complete, we can obtain the agent saying that it cannot eat the apple because it is an agent (or because it is not a human) or other more complex, it would depends only on how much complete is the agents' knowledge and its representation.

Concretely, then, we can change the practical reasoning cycle in Figure 2 by the following actions, which replace lines 10 and 11.

For each action, the agent evaluates the pre- and post-conditions, and then a function we call *rehearsal* is implemented to reason about actions, beliefs, and goals, i.e., desires in the case of BDI agents. With the result of the reflections, the agent updates the knowledge base and then justifies its actions.

In the next section, we give a simple example with a speech act in a well-known scenario developed in the *Jason* online tutorial. This example is used to validate the theoretical approach and verify the use of speech acts for message return and external communication.

foreach α_i **do**
    $evaluate(\alpha_i)$;
    $R \leftarrow rehearsal(\alpha_i, B_{\alpha_i}, D)$;
    $update(B, D)$;
    $J \leftarrow justify(\alpha_i, B_{\alpha_i})$;
end

Algorithm 1: The algorithm for extending the practical reasoning with the inner speech.

7. Implementing Inner Speech with a BDI Example

In the example, a collaborative task between agents is proposed, and no human is present. However, this does not affect the validation, since we want to check here the possibility that an agent activates its inner speech. The example was taken from the examples on *Jason*'s reference page[1].

The scenario, adapted for our purposes, involves collaboration and communication between four agents to complete a construction. The main purpose is to find resources of different types on the map, mine them and bring them to the agent in charge of the construction, called the *Builder*. The Builder agent needs three types of resources to achieve its goal, and works with three Collector agents to do so. Each Collector searches for a resource and can bring it along. Whenever a Collector agent finds a resource that the Builder needs, it picks it up and brings it to him. When a particular type of resource is no longer needed or when the constructions have finished, the builder informs the collectors. The environment is represented through a simple grid representation and collectors are able to perceive whether or not a resource is present in a cell.

The Builder agent specifies the resource to search for, which exists in at most five per type. The Collector agents start searching for the first resource and stop when they have found five. In the original example, this control is centralized to the Builder agent and the Collectors do nothing but execute what is requested by the Builder. To validate the inner language module, we made some changes in this example and moved the resource control to the Collectors. In this way we can insert the *rehearsal* function at the moment the collector agent finds a resource that is not of the type requested by the builder or is no longer needed.

In the following algorithms we show an excerpt of the implemented code for the Collector agent.

The first portion of the code concerns the search for resources and four different cases may occur. In the first two cases, the agent keeps moving since it has not yet found any resources or has found a resource type that is not required. In the second case, the agent finds a resource that it believes is needed (because the collected quantity is not yet sufficient). The simplest case is seen in line 3, where the agent evaluates the triggering event and then changes its beliefs by a speech act directed to itself. The result of the execution can be seen in Figure 5, where the console contains a message that refers to that specific belief.

[1] http://jason.sourceforge.net/wp/examples/, this example was originally written by Rob Clarke and Andy Buck as 2nd coursework for the Multi-Agent Systems Module run in 2004-2005 at the University of Durham, U.K. , later edited by Rafael Bordini.

When the agent is supposed to continue with the search, triggers the next plan (the last case) which involves extracting the found resource and transporting it to the builder. After which it may continue with the extraction or the research.

```
1: +!check_for_resources : not found(_) ← move_to(next_cell).
2: +!check_for_resources : found(R) & enough(R) ← move_to(next_cell).
3: +!check_for_resources : found(R) & not resource_needed(R)
       ← .my_name(Me);
         .send(Me, tell, resource_needed(R));
         !check_for_resources.
4: +!check_for_resources : found(R) & resource_needed(R)
       ← !stop_checking;
         !take(R, builder);
         !continue_mine.
```

Algorithm 2: Part of the code that shows the plans used for searching resources.

```
1: +enough(R)[builder] : found(R) & resource_needed(R)
       ← .my_name(Me);
         .send(Me, untell, [resource_needed(R), found(R)]);
         .drop_all_desires;
         !continue_mine.
2: +enough(R)[builder] : resource_needed(R)
       ← .my_name(Me);
         .send(Me, untell, resource_needed(R)).
```

Algorithm 3: Part of the code that shows the plans that manage the receipt of notification by the builder of a resource type that is no longer needed.

The second part of the code concerns receiving a message from the builder that it has collected enough of a particular resource. In this case, the agent must distinguish between two different circumstances. When it founds a resource of the type specified by the builder, it will perform internal reasoning to remove everything related to that resource from the belief base, except for the new knowledge. It can also drop the task it is performing to continue with the research. Otherwise, it can be enough to update its own belief base.

In this example, we only show the existing mapping between the psychological concept of inner speech and the communication agent module through the use of the speech act. The example presents a complex structure including the environment in which the implemented methods can be interwoven with the rehearsal function. Next step, we will include the *rehearsal R* function for supporting the environment's artifacts and plan library revision.

8. Discussions and Conclusions

Interest in systems capable of self-adaptive and self-aware capabilities is growing rapidly in these years. Equipping robots or agents with cognitive capabilities is certainly the next

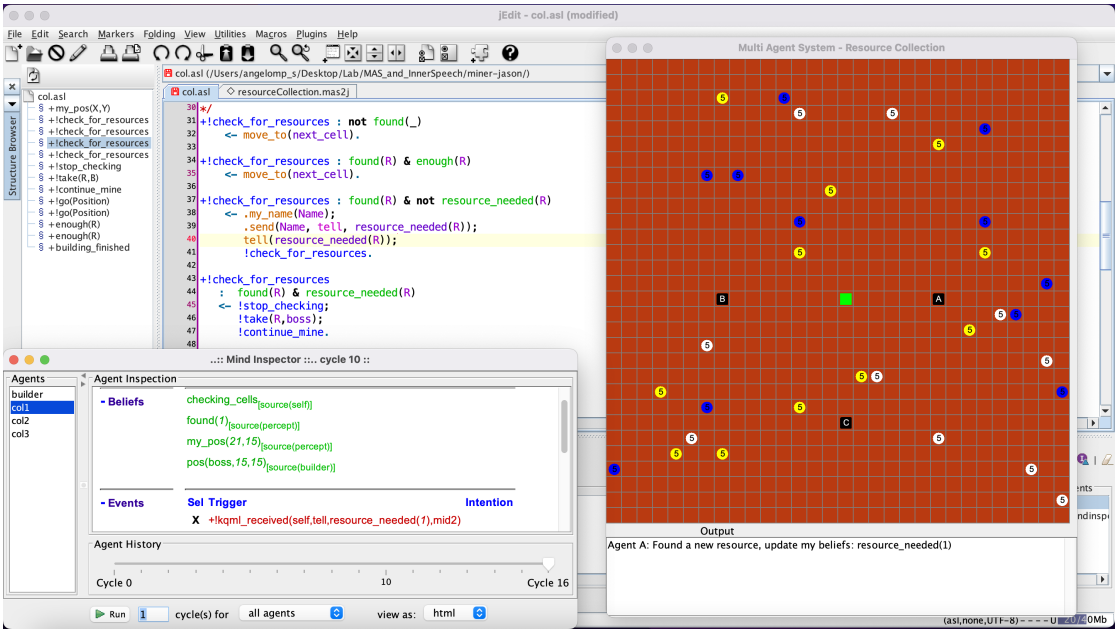

Figure 5: A screenshot of the plans running for the Collector agent.

breakthrough in the field of artificial intelligence, and more and more scientists are talking about machines capable of behaving like humans. In this paper we present a possible solution to endow agents, robots and intelligent systems in general with internal language capabilities. The deliberations one makes before taking an action or making a decision are a key moment for adaptive and autonomous behavior, especially when working in teams. Humans have developed the ability to put themselves at the center of their thinking and to activate what is known as inner discourse to regulate and control their behavior. The idea we present in this article is an preliminary approach to use the concept of speech act to implement inner speech in agents. From a technological point of view, in moving from theory to implementation, we experimented with the agent's BDI technology and obtained good results in terms of simplicity of handling some specific design abstractions. We then proposed a validation of the theoretical approach through a simple example of a multi-agent system developed in *Jason*.

The proposed approach builds on and extends our previous work on endowing robots with the ability to justify the results of their actions. The ability to explain what is done and why is the focus of our work, which aims to create agents that are reliable, explainable, and believable. We believe that the method we use and the cognitive model that underlies all of our work can also incorporate other elements such as emotions, mental states in general, and even moral and ethical values into the reasoning process. This will be the topic of our future work.

The advantage we drew from the choices we made was that we could easily connect the theoretical model to its practical implementation by using the *Jason* interpreter's reasoning cycle. However, a challenge that we highlighted during our work, and that will be the subject of our future work, is that the transition from the cognitive model to the implemented model requires a careful and precise methodological approach to the design. The high-level requirements that

we have highlighted in this paper, especially those that need to be considered for the runtime phase (belief, action, goal, plan), need to be carefully analysed and transformed into the typical *Jason* implementation elements, especially when scenarios are complex.

Another element to be considered is the completeness and accuracy of the agent's knowledge representation. Indeed, it is necessary that the elements of the environment are correctly associated with all the actions that can be performed and the rules for their activation. Ontologies, which can be found in the literature and for which a process of conversion to belief is necessary, can help in this regard, and their integration will be part of our future work.

Finally, we plan to further validate the proposed approach on a more complex scenario, e.g., the table setting, and complement it with a rigorous methodological approach.

9. Acknowledgments

The research leading to these results is in the frame of the project awarded by number FA9550-19-1-7025 - Air Force Office of Scientific Research.

References

[1] A. S. Rao, M. P. Georgeff, et al., BDI agents: from theory to practice., in: ICMAS, volume 95, 1995, pp. 312–319.

[2] A. S. Rao, Agentspeak (l): BDI agents speak out in a logical computable language, in: European Workshop on Modelling Autonomous Agents in a Multi-Agent World, Springer, 1996, pp. 42–55.

[3] M. Georgeff, A. Rao, Rational software agents: from theory to practice, in: Agent technology, Springer, 1998, pp. 139–160.

[4] R. H. Bordini, J. F. Hübner, BDI agent programming in agentspeak using jason, in: International Workshop on Computational Logic in Multi-Agent Systems, Springer, 2005, pp. 143–164.

[5] O. Boissier, R. H. Bordini, J. F. Hübner, A. Ricci, A. Santi, Multi-agent oriented programming with jacamo, Science of Computer Programming 78 (2013) 747–761.

[6] A. Chella, F. Lanza, V. Seidita, A cognitive architecture for human-robot teaming interaction, in: Proceedings of the 6th International Workshop on Artificial Intelligence and Cognition, Palermo, 2018.

[7] J. E. Laird, C. Lebiere, P. S. Rosenbloom, A standard model of the mind: Toward a common computational framework across artificial intelligence, cognitive science, neuroscience, and robotics, Ai Magazine 38 (2017) 13–26.

[8] J. R. Anderson, The architecture of cognition, volume 5, Psychology Press, 1996.

[9] A. Chella, F. Lanza, V. Seidita, Human-agent interaction, the system level using jason, in: Proceedings of the 6th International Workshop on Engineering Multi-Agent Systems (EMAS 2018), Stockholm, 2018.

[10] V. Seidita, C. Diliberto, P. Zanardi, A. Chella, F. Lanza, Inside the robot's mind during human-robot interaction, in: 7th International Workshop on Artificial Intelligence and Cognition, AIC 2019, volume 2483, CEUR-WS, 2019, pp. 54–67.

[11] C. Castelfranchi, A. Chella, R. Falcone, F. Lanza, V. Seidita, Endowing robots with self-modeling abilities for trustful human-robot interactions, in: 20th Workshop" From Objects to Agents", WOA 2019, volume 2404, CEUR-WS, 2019, pp. 22–28.
[12] C. Castelfranchi, R. Falcone, Trust theory: A socio-cognitive and computational model, volume 18, John Wiley & Sons, 2010.
[13] R. Falcone, C. Castelfranchi, Trust dynamics: How trust is influenced by direct experiences and by trust itself, in: Autonomous Agents and Multiagent Systems, 2004. AAMAS 2004. Proceedings of the Third International Joint Conference on, IEEE, 2004, pp. 740–747.
[14] A. Morin, A neurocognitive and socioecological model of self-awareness, Genetic, social, and general psychology monographs 130 (2004) 197–224.
[15] P. J. Silvia, T. S. Duval, Objective self-awareness theory: Recent progress and enduring problems, Personality and social psychology review 5 (2001) 230–241.
[16] S. Duval, R. A. Wicklund, A theory of objective self awareness. (1972).
[17] A. Morin, Possible links between self-awareness and inner speech theoretical background, underlying mechanisms, and empirical evidence, Journal of Consciousness Studies 12 (2005) 115–134.
[18] L. E. Berk, R. A. Garvin, Development of private speech among low-income appalachian children., Developmental psychology 20 (1984) 271.
[19] A. E. Winsler, C. E. Fernyhough, I. E. Montero, Private speech, executive functioning, and the development of verbal self-regulation., Cambridge University Press, 2009.
[20] L. S. Vygotsky, M. Cole, Mind in society: Development of higher psychological processes, Harvard university press, 1978.
[21] J. R. Searle, J. R. Searle, Speech acts: An essay in the philosophy of language, volume 626, Cambridge university press, 1969.
[22] J. L. Austin, How to do things with words, Oxford university press, 1975.
[23] R. H. Bordini, J. F. Hübner, M. Wooldridge, Programming multi-agent systems in AgentSpeak using Jason, volume 8, John Wiley & Sons, 2007.
[24] G. K. Soon, C. K. On, P. Anthony, A. R. Hamdan, A review on agent communication language, Computational Science and Technology (2019) 481–491.
[25] T. Finin, R. Fritzson, D. McKay, R. McEntire, Kqml as an agent communication language, in: Proceedings of the third international conference on Information and knowledge management, 1994, pp. 456–463.
[26] T. Finin, R. Fritzson, D. P. McKay, R. McEntire, et al., Kqml-a language and protocol for knowledge and information exchange, in: 13th Int. Distributed Artificial Intelligence Workshop, 1994, pp. 93–103.
[27] P. D. O'Brien, R. C. Nicol, Fipa—towards a standard for software agents, BT Technology Journal 16 (1998) 51–59.
[28] M. T. Kone, A. Shimazu, T. Nakajima, The state of the art in agent communication languages, Knowledge and Information Systems 2 (2000) 259–284.
[29] F. Bellifemine, A. Poggi, G. Rimassa, Developing multi-agent systems with a fipa-compliant agent framework, Software: Practice and Experience 31 (2001) 103–128.
[30] Y. Labrou, T. Finin, A semantics approach for kqml—a general purpose communication language for software agents, in: Proceedings of the third international conference on Information and knowledge management, 1994, pp. 447–455.

[31] J. Mayfield, Y. Labrou, T. Finin, Evaluation of kqml as an agent communication language, in: International Workshop on Agent Theories, Architectures, and Languages, Springer, 1995, pp. 347–360.

[32] A. Chella, F. Lanza, V. Seidita, Decision Process in Human-Agent Interaction: Extending Jason Reasoning Cycle, in: Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics), volume 11375 LNAI, Springer, Cham, 2019, pp. 320–339. doi:10.1007/978-3-030-25693-7_17.

[33] A. Chella, F. Lanza, A. Pipitone, V. Seidita, Knowledge acquisition through introspection in Human-Robot Cooperation, Biologically Inspired Cognitive Architectures 25 (2018) 1–7. doi:10.1016/j.bica.2018.07.016.

[34] A. Chella, F. Lanza, V. Seidita, Representing and developing knowledge using Jason, CArtAgO and OWL, in: CEUR Workshop Proceedings, volume 2215, 2018, pp. 147–152.